



L O V E L Y
P R O F E S S I O N A L
U N I V E R S I T Y

CSB205: SYSTEM DESIGN Major Project

Topic: Multi-Tenant E-Commerce Marketplace with Search & Recommendations

Team Names: Perineni Lakshith, Rachana Reddy, Jashwanth Raparthi, Praveen Kumar, Rakesh Reddy

Team Reg.No's: 12401218, 12400989, 12406316, 12405070, 12406880

Section: K24BX

1. Requirements Pack

1.1 Stakeholders

A successful marketplace requires satisfying diverse needs across three main user groups and the core technology team.

Stakeholder Group	Core Needs / Domain Focus
Sellers (Tenants)	Storefront customization, inventory management, order processing, sales analytics, billing.
Buyers (End Users)	Personalized product search, seamless cart/checkout, order tracking, reviews, secure payments.
Platform Administrators	Tenant oversight, policy enforcement, platform-wide configurations, system health monitoring.
Development/DevOps Team	CI/CD automation, easy scaling, robust monitoring, fault isolation, cost efficiency.

1.2 User Stories (Feature Context)

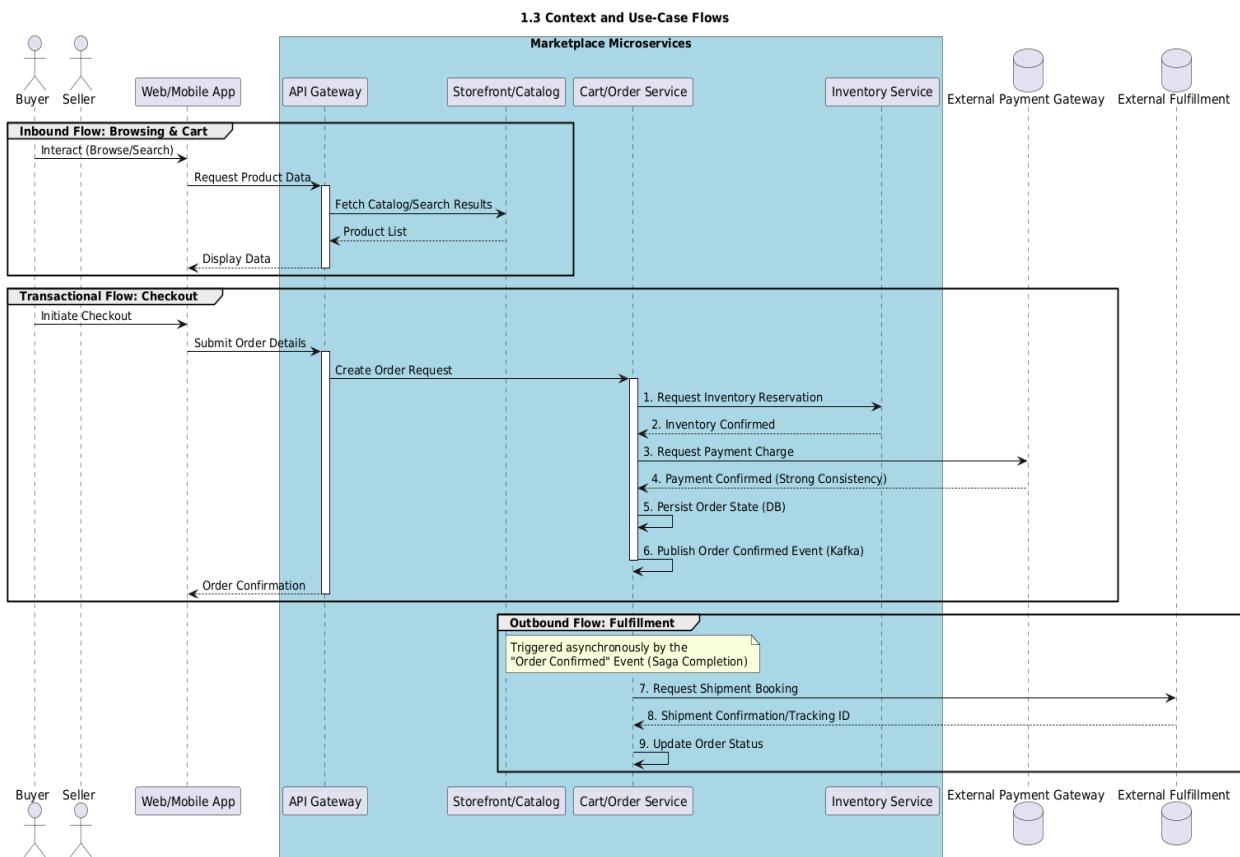
Persona	User Story	Core Feature Implied
Seller	"As a Seller, I want to upload bulk inventory via CSV so that I can quickly populate my catalog."	Catalog & Inventory Service, Data Ingestion Pipeline
Buyer	"As a Buyer, I want search results to prioritize items based on my browsing history so that I find relevant products faster."	Search Service, Recommendation Engine
Admin	"As an Admin, I want to disable a tenant's storefront immediately if"	Tenant Service, Identity & Access Control

	they violate policy."	
--	-----------------------	--

1.3 Context and Use-Case Diagrams

The system operates as a central nervous system coordinating interactions between end-users, tenants, and various external systems.

- **Inbound Flow:** Buyers interact via Web/Mobile App -> API Gateway -> Storefront Service (for browsing) or Cart/Order Service (for transactions).
- **Outbound Flow:** Order Service -> Payment Gateway (e.g., Stripe, PayPal) -> Shipping/Fulfillment Provider (e.g., FedEx, UPS).



1.4 Constraints and Assumptions

Category	Constraint	Assumption
Non-Functional	Elastic scale for traffic spikes, strong consistency for payments/orders, strict PII protection.	APIs (Payments, Fulfillment) are generally reliable and available.
Technical	Architecture must use Domain-Driven Design (DDD) principles and leverage event choreography.	The team is proficient in multiple languages (Go/Java) and Kubernetes.
Regulatory	PCI-aware design for payment handling, compliance with data privacy laws (GDPR/CCPA).	PCI scope is minimized via external payment processors (tokenization).

2. Architecture

2.1 Trade-off Analysis of Styles

The **Microservices Architecture** was chosen to meet the non-functional goals of elastic scale, fault isolation, and technology heterogeneity required by a modern, high-traffic marketplace.

Style	Rationale for Rejection/Acceptance	Trade-offs Accepted
Monolith	Rejected. Fails to meet elastic scale requirement (e.g., scaling Search independently). Leads to a single point of failure.	N/A
Microservices	Selected. Allows independent deployment,	Accepted: Operational complexity, distributed

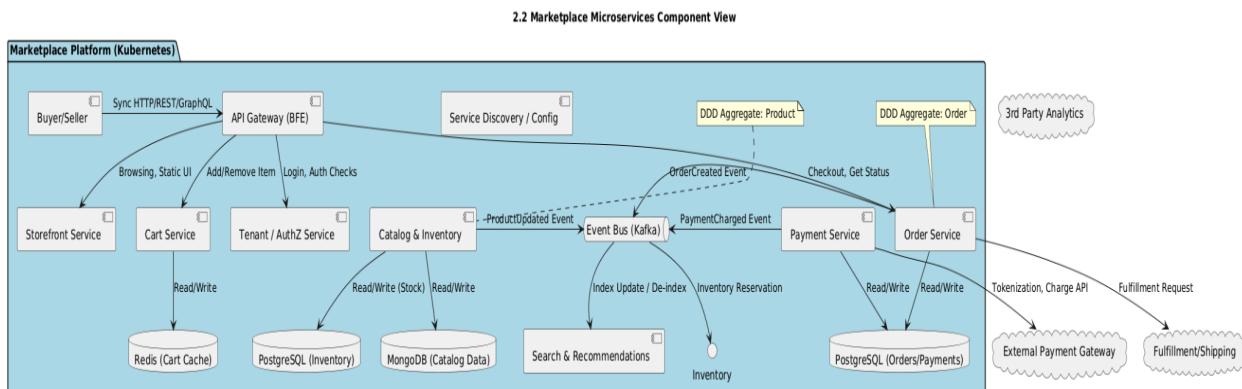
	technology specialization (e.g., Go for performance, Java for complex logic), and fault boundaries.	transaction handling, eventual consistency.
Serverless	Hybrid Use. Suitable for asynchronous tasks (e.g., image resizing, report generation) but poor for long-running transactional Sagas.	Accepted: Potential cold start latency for non-critical paths.

2.2 Component Diagram

The system is divided into core domains, communicating asynchronously via an Event Bus (Kafka) and synchronously via the API Gateway.

Key Services/Domains and DDD Contexts:

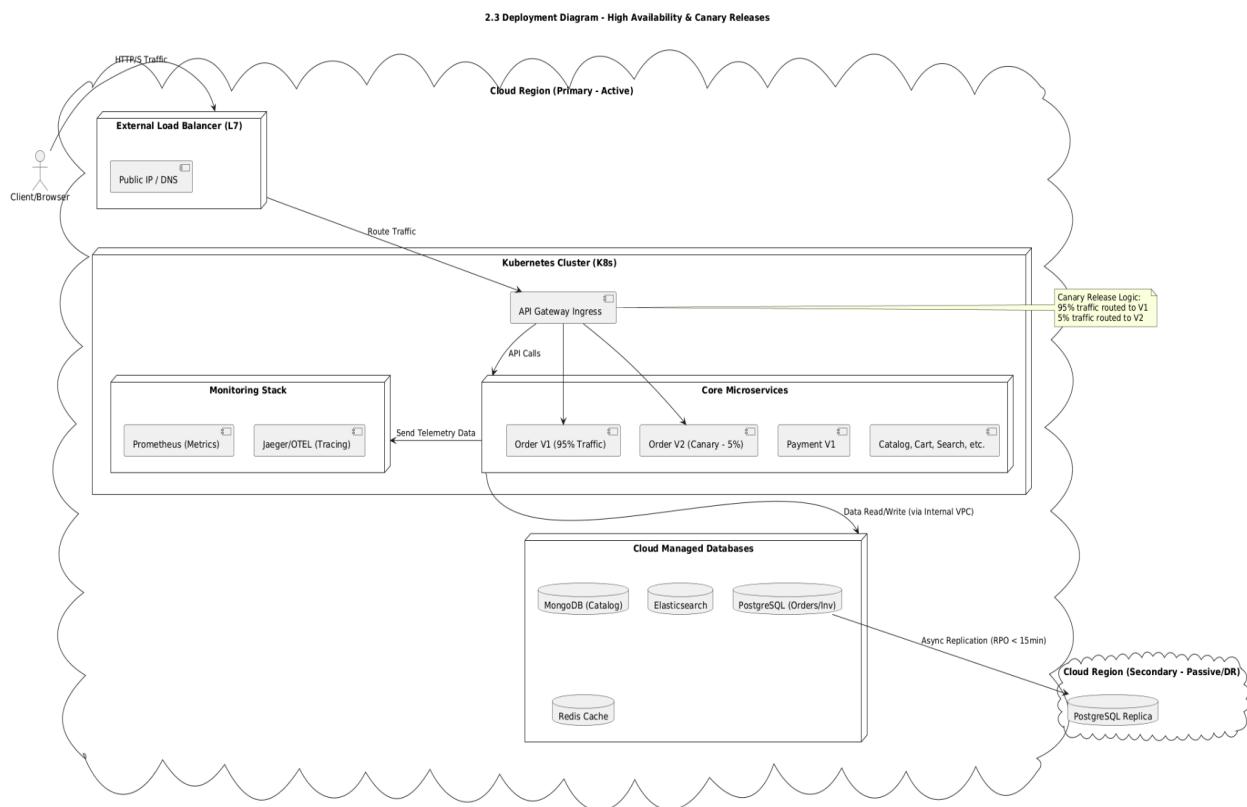
- **API Gateway (BFE):** Acts as a Backend-for-Frontend, responsible for request aggregation (e.g., fetching product, reviews, and inventory in one call for a Product Detail Page), authentication, and rate limiting.
- **Storefront:** Buyer UI/UX, static content delivery.
- **Catalog & Inventory:** Product details (NoSQL), stock levels (Relational). **DDD Detail:** The Product is the Aggregate Root here.
- **Search & Recommendations:** Elasticsearch/AI components for personalization.
- **Cart:** Manages buyer sessions, uses Redis cache.
- **Order:** Manages the order lifecycle (DDD Aggregate Root). All state changes are validated within this boundary.
- **Payment:** Handles tokenization, reconciliation, and provider communication.
- **Tenant/AuthZ:** Manages seller accounts, roles, and permissions.



2.3 Deployment Diagram

The architecture is deployed on Kubernetes (K8s) for orchestration, enabling high availability and controlled releases.

- **Platform:** Kubernetes (K8s) for container orchestration, providing self-healing and service discovery.
- **Load Balancing:** External Load Balancer (L7) routes traffic to the API Gateway.
- **Deployment Strategy (Canary Releases):** For critical services (e.g., Payment, Order), K8s is configured to use **Canary Releases**. New versions are initially rolled out to a small subset (e.g., 5%) of production traffic, monitored extensively for performance and error rates via Prometheus/Jaeger, before proceeding with a full deployment roll-out. This minimizes blast radius during software updates.
- **Database:** Cloud-managed services (e.g., AlloyDB/PostgreSQL, Managed MongoDB/CosmosDB, Elasticsearch Cluster) to reduce operational overhead.



3. Design Patterns

The architecture relies heavily on domain-specific patterns to manage complexity and distributed state.

Pattern	Type	Implementation Detail	Anti-Patterns Avoided
Saga	Behavioral	Order Lifecycle: Uses event choreography (e.g., OrderCreated -> InventoryReserved -> PaymentCharged). Guarantees eventual consistency for distributed transactions.	Two-Phase Commit (2PC) in a distributed system (too slow/complex).
Outbox	Behavioral	Ensures that database updates (e.g., Order status change) and the publishing of corresponding events (e.g., OrderConfirmed) are atomic via a dedicated outbox table.	Dual write problem (inconsistent state between DB and Event Bus).
Factory	Creational	Payment Provider Factory: Decouples the Payment Service from concrete payment gateways (Stripe, PayPal) by creating provider-specific	Excessive conditional logic (if/else hell) for external system integration.

		classes via a common interface.	
Adapter	Structural	Fulfillment Adapter: Standardizes the interface for heterogeneous external fulfillment services (FedEx, UPS) so the Order Service doesn't rely on vendor-specific logic.	Direct reliance on external vendor APIs, making vendor switching difficult.
Facade	Structural	Search Facade: Provides a simple search(query) interface to the Storefront, hiding the complexity of the Elasticsearch query language, ranking, and personalization logic.	Clients needing deep knowledge of the underlying search technology.

4. Database Design

4.1 Polyglot Persistence Rationale

Choosing the right data store per service optimizes performance and consistency for specific domain needs.

- **Relational (PostgreSQL/AlloyDB):** Used for **Orders, Payments, User Accounts.**
Rationale: Non-negotiable strong consistency (ACID), transactional integrity, and complex financial reporting where joins are necessary.
- **Document Store (MongoDB):** Used for **Product Catalog.** *Rationale:* Flexible schema to support diverse product attributes (e.g., clothes vs. electronics) across different sellers (tenants) without schema migrations.

- **Search Index (Elasticsearch):** Used for **Product Discovery, Recommendations.**
Rationale: Optimized for full-text search, fuzzy matching, and fast aggregation (faceting) required by the Buyer experience.
- **Cache (Redis):** Used for **Cart Data, Session State, Popular Product Lookup.**
Rationale: High-speed, low-latency key-value store, supporting ephemeral state.

4.2 Entity-Relationship Diagram (ERD) / Schema Overview

4.3 Normalization/Denormalization Rationale

- **Normalization (3NF):** Applied strictly to **Orders/Payments** tables within Postgres. Ensures data integrity, prevents update anomalies, and supports complex joins for financial audits.
- **Denormalization:** Applied in the **Catalog Service** (embedding product features, reviews summaries) and heavily in the **Search Index** (flattening product data for fast read access). *Rationale:* Optimized for read-heavy operations, trading write complexity for massive read speed improvements.

4.4 Indexing, Partitioning/Sharding Plan

- **Indexing:** All relational tables must be indexed on the composite key (`tenant_id`, `primary_key`) for efficient multi-tenant query isolation. Secondary indexes on (`tenant_id`, `created_at`) are critical for tenant-scoped reporting.
 - **Partitioning/Sharding (Horizontal Scaling):**
 - **Strategy:** Shard data primarily by `tenant_id`. This achieves **Tenant Isolation** and distributes the load from large tenants across different database instances, minimizing the "noisy neighbor" effect.
 - **High-Volume Tables (Orders):** Further partitioning by date range (e.g., monthly) within a shard can improve archive and query performance.
 - **Data Synchronization (CDC):** Data is synchronized between the transactional (Postgres/Mongo) and search (Elasticsearch) stores using the **Change Data Capture (CDC)** mechanism. The Outbox pattern ensures every relevant database change generates an event, which is then consumed by the Search Service to maintain the up-to-date index, guaranteeing read-model eventual consistency.
-

5. Performance & Scale

5.1 Caching (In-Memory/CDN)

Cache Layer	Technology	Data Stored	Strategy
L3 (CDN)	Cloudflare/Akamai	Static assets, Seller logos, Product images.	Edge caching, Geo-distribution for global user base.
L2 (Distributed)	Redis Cluster	Shopping Carts, Session State, Personalized Recommendations, Hot Product IDs.	TTL-based expiration, High-speed lookup, handles sudden spikes in read traffic.
L1 (In-Memory)	Service Local Cache	Feature flags, Configuration values, Tenant account details (frequently accessed).	Short TTL, small data sets, eliminates network hop for critical config.

5.2 Load Balancing Strategy

- **Layer 7 (Application Load Balancer - ALB):** Used at the API Gateway level. It can inspect HTTP headers and hostnames, allowing for intelligent routing based on tenant subdomain or path.
- **Strategy:** Primarily **Least Connections** to ensure new requests are routed to the least-busy service instance, maximizing utilization and minimizing individual instance load.

5.3 Replication

- **Database:** Asynchronous streaming replication (read replicas) for all relational stores (Postgres) to handle heavy read traffic (e.g., reports, catalog views) and ensure database resilience.
- **Services:** All microservices run with a minimum of 3 replicas across multiple Availability Zones (AZs) in K8s to withstand zone-level failures.

5.4 Backpressure

- **Mechanism:** Implemented at the API Gateway using the **Token Bucket Algorithm**.
 - **Function:** Enforces **rate limiting** policies configured per tenant_id and global service limits. This prevents a surge of traffic from one misbehaving tenant (a "noisy neighbor") from degrading service quality for all others, ensuring system stability.
-

6. Security & Reliability

6.1 Threat Model (Core Defenses)

Threat Category	Example Threat	Primary Defense Mechanism
Access Control	Unauthorized tenant accessing another tenant's data.	Row-Level Security (RLS) in databases, enforced checks in services using tenant_id.
Denial of Service (DoS)	Traffic flooding the search endpoint.	WAF (Web Application Firewall) for L7 attacks, and Rate Limiting at the API Gateway.
Data Exposure	Payment details interception.	PCI-aware design (tokenization), TLS 1.3 encryption end-to-end, Secret Rotation (using KMS/Vault).

6.2 Authentication (AuthN) and Authorization (AuthZ)

- **AuthN:** Handled by a centralized **Identity Provider (IdP)** (e.g., Keycloak). Uses **OIDC** and **JWT (JSON Web Tokens)** for stateless session management.
- **AuthZ: Role-Based Access Control (RBAC)** enforced at two critical levels:
 1. **API Gateway:** Validates JWT structure and scopes/roles before routing.
 2. **Service Level:** Each service validates the user's tenant_id and role against the

resource being accessed.

- **Tenant Isolation via RLS:** Row-Level Security (RLS) is paramount for multi-tenancy. In PostgreSQL, RLS policies are enforced to ensure that a query originating from a service authenticated with `tenant_id=X` can *only* retrieve or modify rows where the `tenant_id` column equals X. This provides a fundamental security layer against cross-tenant data access attempts.

6.3 OWASP Top 10 Defenses

- **A03: Injection:** Use **Parameterized Queries/ORMs** (e.g., Hibernate, GORM) instead of string concatenation for database access to eliminate SQL Injection risk.
- **A04: Insecure Design/Broken Access Control:** Enforce multi-layer access checks, especially using the `tenant_id` filter on all read/write operations.
- **A05: Security Misconfiguration:** Use infrastructure-as-code (IaC) and secret rotation services to manage configuration securely.

6.4 Resilience Patterns

- **Circuit Breakers:** Implemented using libraries (e.g., Hystrix, Resilience4j) on every synchronous inter-service call. If the Review Service fails, the Order Service opens the circuit, and the Product Page loads without reviews rather than crashing or freezing (graceful degradation).
- **Retries:** Use exponential backoff with jitter for transient errors (e.g., network timeout) before failing the transaction.

6.5 Disaster Recovery (DR) Plan

- **Strategy:** Active-Passive (Warm Standby) Multi-region deployment. The Passive region maintains running instances and up-to-date data replicas.
 - **RPO (Recovery Point Objective):** > 15 minutes (Achieved via asynchronous database replication and frequent hourly backups).
 - **RTO (Recovery Time Objective):** > 1 hour (Achieved via automated K8s failover and DNS switchover scripts, minimizing downtime).
-

7. API Specification

7.1 Protocol Selection

- **REST (Synchronous):** Used for **Admin and Transactional APIs** (/orders, /payments, /inventory). Standard CRUD operations.
- **GraphQL/gRPC (Synchronous):** **GraphQL** is used for the **Storefront API** to allow clients (Web/Mobile) to fetch complex, nested data (Product, Reviews, Related Items) in a single optimized query.
- **Events (Asynchronous):** Used for **Order Orchestration** (OrderCreated, InventoryReserved) via Kafka.

7.2 Endpoints (REST Example)

Path	Method	Purpose	Idempotency	Error Codes
/api/v1/orders	POST	Create a new order.	Key Required	400, 403, 409 (Conflict - if key exists)
/api/v1/orders/{id}	GET	Retrieve specific order details.	N/A	404, 403
/api/v1/payments	POST	Process payment for a transaction.	Key Required	402 (Payment Failed)

7.3 Versioning

- **Strategy: URI Versioning** (/api/v1/orders).
 - **Rationale:** Provides clear, non-breaking contracts to external consumers and simplifies parallel deployment for new major feature releases.
-

8. Observability

Observability is built upon the three pillars to ensure rapid detection and diagnosis of issues in the distributed environment.

Pillar	Technology/Tool	Data Collected	Use Case
Logs	ELK Stack (Elasticsearch, Logstash, Kibana) / Datadog	Structured JSON logs from all services.	Root cause analysis, detailed debugging of error states.
Metrics	Prometheus & Grafana	RED (Rate, Errors, Duration) metrics for every service, CPU/Memory utilization, Queue depth.	Health monitoring, capacity planning, SLO tracking.
Traces	OpenTelemetry / Jaeger	Trace ID and Span context propagation across all service calls.	Identifying latency bottlenecks in distributed Sagas across multiple services.

8.1 SLOs/SLIs and Alerting Runbook

- **SLI (Service Level Indicators):**
 - **Availability:** Success rate of critical API calls (e.g., POST /orders success rate \$> 99.91%).
 - **Latency:** P95/P99 latency of API calls (e.g., P99 of /search must be \$le 200\$ms).
 - **SLO (Service Level Objective):** Critical service availability target set at **99.9%**.
 - **Alerting Runbook:** Automated alerts (e.g., PagerDuty) triggered when SLI breaches SLO. Each alert contains a link to a runbook detailing immediate triage steps (e.g., check database replication lag, scale up Cart Service replicas, analyze Kafka consumer lag).
-

9. Tech Stack Justification

Tier	Selected Option(s)	Alternatives Considered	Trade-offs / Justification
Backend	Go, Java/Spring Boot	Node.js, Python	Go for high-throughput, low-latency services (Checkout, Cart). Java/Spring for complex, long-running business logic (Order Orchestration, Admin) due to its maturity and strong enterprise features.
Frontend	React	Vue.js, Angular	Rich ecosystem, strong component model, largest pool of experienced developers for rapid development.
Data/Search	Postgres, Mongo, Elasticsearch, Redis	MySQL, Couchbase, Cassandra	Polyglot Persistence: Optimal performance/consistency profile achieved by matching the right tool to the domain need.
Messaging (Event Fabric)	Kafka	RabbitMQ, ActiveMQ	High throughput, durability, persistent event log required for event choreography and replayability

			(Sagas), making it the central nervous system.
Infrastructure	Kubernetes (K8s)	Docker Swarm, ECS	Industry standard for microservices orchestration, provides advanced networking, auto-scaling, and operational resilience features, reducing vendor lock-in.